

# Implementing Temporal Features in PostgreSQL: SQL Standard and Beyond

Hettie Dombrovskaya DRW  
Boris Novikov

# Who am I

Hettie Dombrovskaya

Database Architect at DRW Holdings Chicago IL

Local Organizer of Chicago PostgreSQL User Group



# Temporal DB Basics



## Outline

- Time travel
- Cloud
- Warehouses
- User-defined functions (UDFs)



## Why Time Travel?

- ✓ Point-in-time (snapshot) queries: How my report looked on the last day of previous month?
- ✓ Change Log: When and how the state of my request was changed?
- ✓ Fully Temporal:
  - ☐ When these objects co-exist?
  - ☐ Before/after/meets etc.
  - ☐ Temporal joins, aggregations etc.



## Potential Solutions

- ✓ Snapshots: any theoretical paper starts with it
- ✓ Event Logs: Often produced by applications with a hope that they will be used for analysis in the future
- ✓ Periods: Used by everyone who actually implements temporal features in a database

***All are equivalent but query performance may differ***





## Time Dimensions

This is something beyond common sense in the real life, however, we need them

- Transactional (system)
- Valid = Effective
- Asserted ~ transactional

Much more were defined but are not widely known

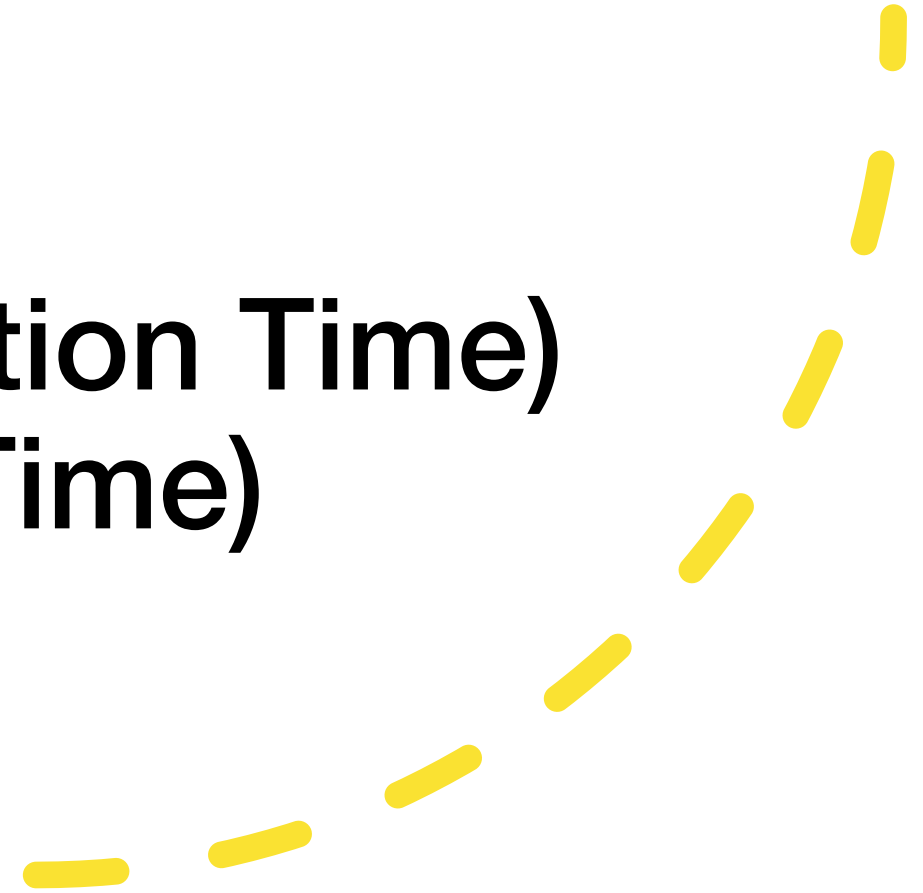
# Temporal Features in the SQL Standard





## Types and Predicates

### The Standard:

- ✓ Does not introduce a period type, instead, a pair of timestamp columns can represent the period
  - ✓ Defines period as closed-open
  - ✓ Supports period predicates: OVERLAPS, MEETS, etc. (Similar but not equivalent to Allen operators)
  - ✓ Supports System time (Transaction Time) and/or Application Time (Valid Time)
- 

## System (Transaction) Time

- System-versioned tables, the name `SYSTEM_TIME` is fixed.
- Before and now, never in the future
- Can never be changed
- Does NOT require separate table for historical data, although some implementations do that
- Default – CURRENT record

## **Application (Valid) Time**

- Maintained by the user
- Column names can be arbitrary
- No semantics are specified
- Currently at most one additional time dimension can be specified

# Primary Keys

The Standard does not provide clean resolution for the PK


ENo	EStart	EEnd	EDept
22217	2010-01-01	2011-09-10	3
22217	2010-02-03	2011-11-12	4

Triple (ENo, EStart, EEnd) does not work

Instead, it suggests “no overlap”

```
ALTER TABLE Emp
```

```
ADD PRIMARY KEY (ENo,  
EPeriod WITHOUT OVERLAPS)
```




# Foreign Keys

Referential integrity constraints should be time-aware, the example below won't work:

ENo	EStart	EEnd	EDept
22217	2010-01-01	2011-09-10	3
22217	2010-02-03	2011-11-12	4

DNo	DStart	DEnd	DName
3	2009-01-01	2011-12-31	Test
4	2011-06-01	2011-12-31	QA

```
ALTER TABLE Emp
  ADD FOREIGN KEY
    (Edept, PERIOD EPeriod)
REFERENCES Dept
  (DNo, PERIOD DPeriod)
```



# Queries and DML

## Syntax extensions for INSERT, UPDATE, DELETE to specify period(s)

```
UPDATE Emp
  FOR PORTION OF EPeriod
    FROM DATE '2011-02-03'
    TO DATE '2011-09-10'
  SET EDept = 4
  WHERE ENo = 22217
```

## Syntax extensions for SELECT

```
SELECT Ename, Edept
FROM Emp
WHERE ENo = 22217 AND
      EPeriod CONTAINS DATE '2011-01-02'
```

# PostgreSQL Support for Temporal Data



## What is Available

- ✓ Period data types are provided although are not required by the Standard
- ✓ Rich set of operators and functions for timestamps, intervals, and periods
- ✓ Predicates are implemented as required by the Standard (including closed-open semantics)

Additionally:

- ✓ GIST indexes
- ✓ GIST with exclusion constraints (solve temporal PK problem)

# What is Missing

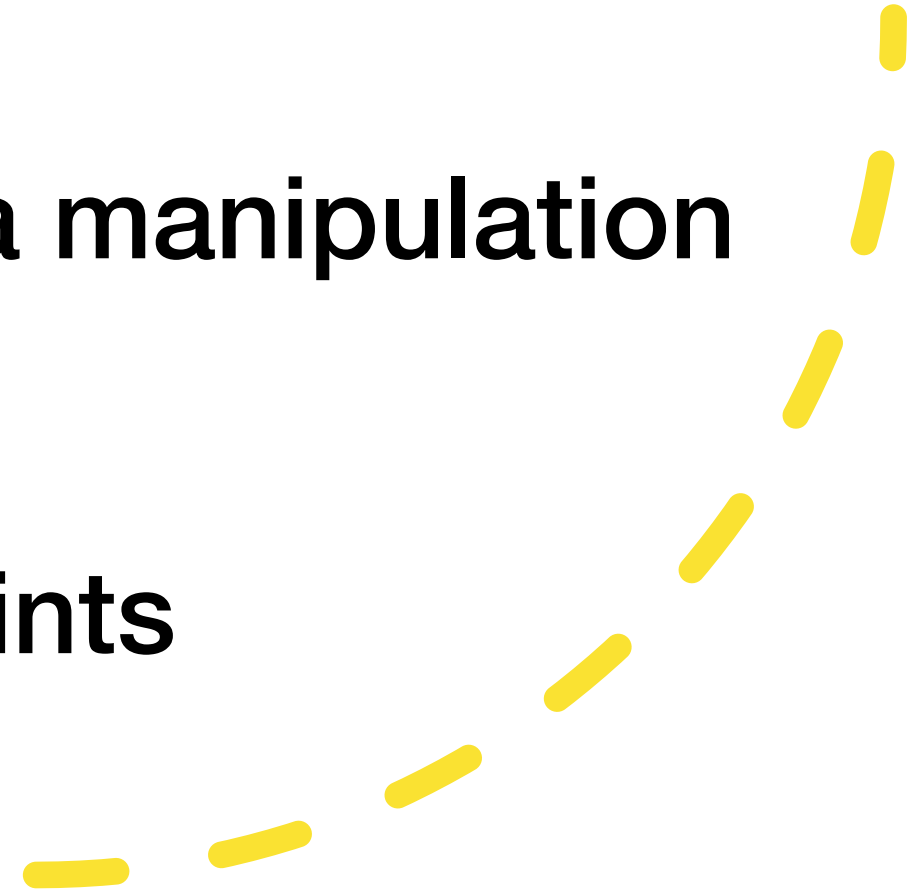
## Language extensions

- ✓ **CREATE temporal table**
- ✓ **SELECT within time period (except for adding explicit condition)**
- ✓ **Modified INSERT/UPDATE/DELETE syntax**

# Bitemporal Implementation

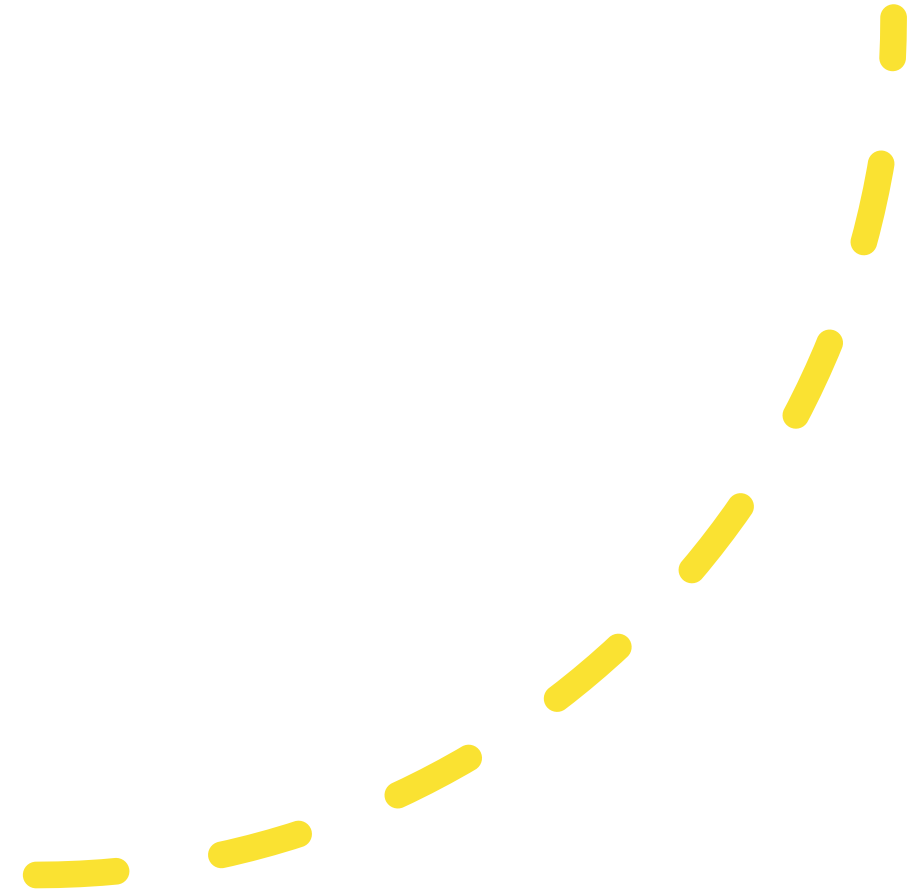


## Overview of Bitemporal Model

- ❑ Supports asserted and effective time dimensions
  - ❑ Transactional dimension can be derived from `row_created_at` timestamp, but is not explicitly supported
  - ❑ Heavily relies on PostgreSQL features
  - ❑ Does not provide any syntax extensions
  - ❑ Data manipulation is implemented with user-defined functions
  - ❑ Provides detailed refinement of data manipulation semantics
  - ❑ Supports temporal integrity constraints
- 

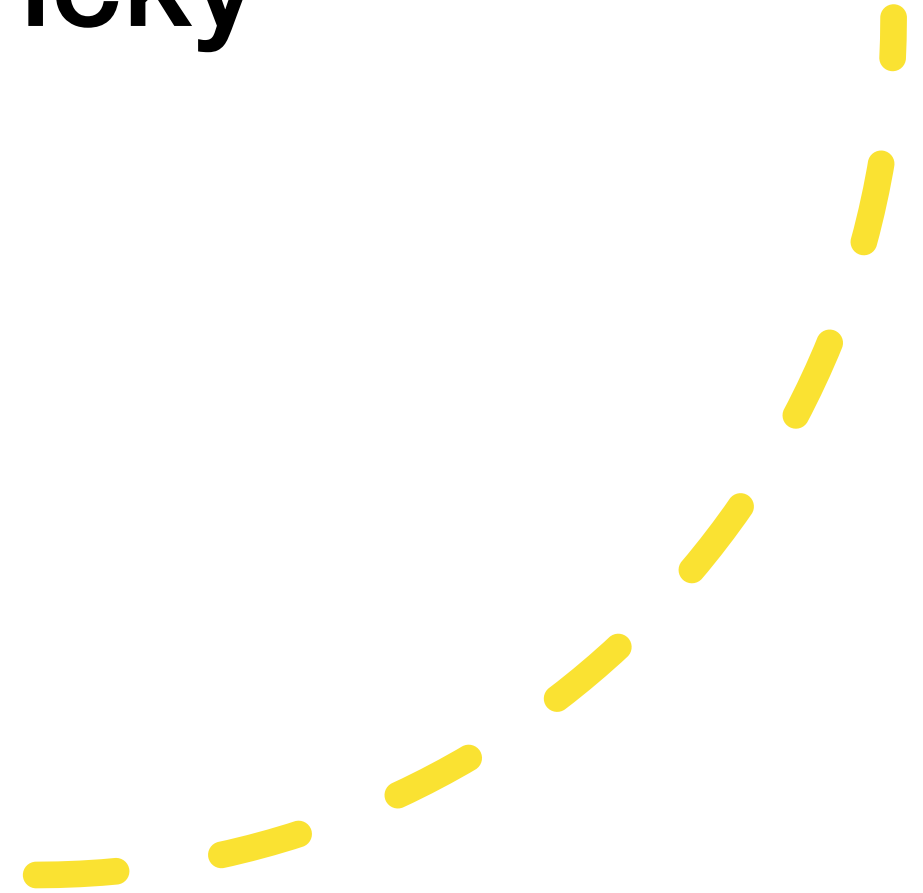


## Storage and Indexes

- ❑ Bitemporal primary keys (business keys) are defined as keys with **NO OVERLAP** utilizing GIST with exclusion constraints.
  - ❑ GIST indexes make bitemporal search efficient
- 



# Queries


- ☐ Time-related conditions must be explicitly specified (acceptable in the Standard)
  - ☐ Built-in predicates (INCLUDES, OVERLAPS, etc.) are helpful
  - ☐ Fully temporal queries are still tricky
- 



# Data Manipulation

## The refinements of manipulation semantics

- ☐ INSERT
- ☐ UPDATE
- ☐ CORRECTION
- ☐ INACTIVATE
- ☐ DELETE





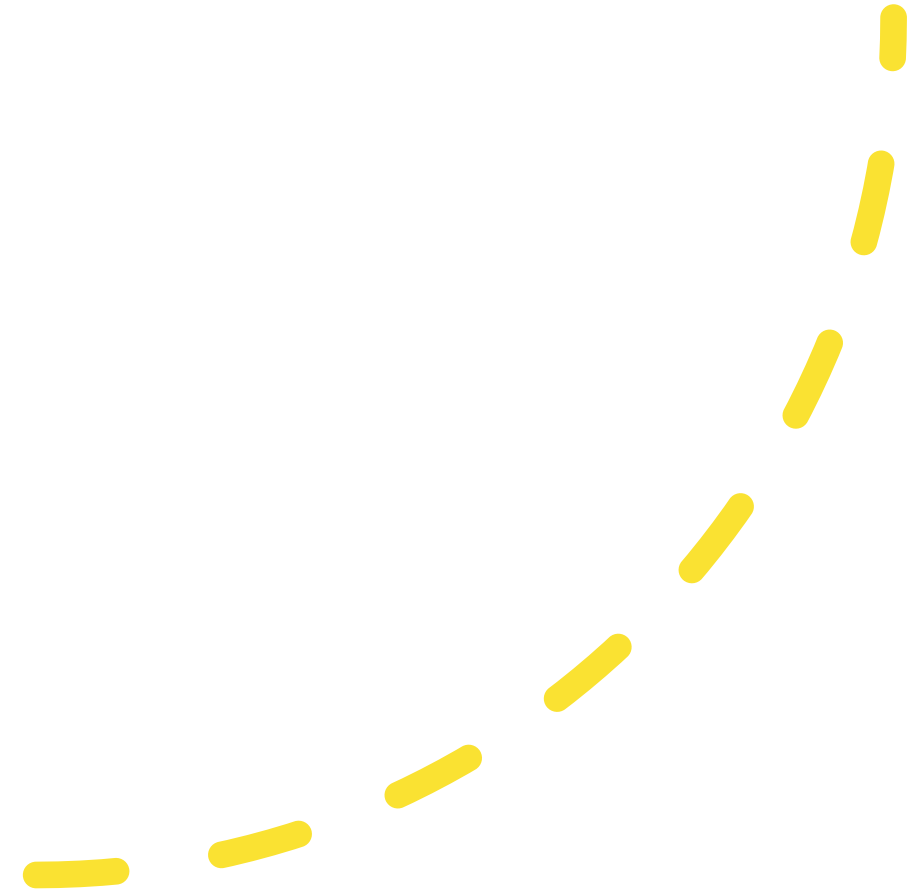
# Bitemporal vs. SQL Standard Comparison

Feature	Standard	Bitemporal
Period type	Not required	✓
Open/close semantics and predicates	✓	✓
SYSTEM_TIME	✓	Implicit
APPLICATION_TIME	✓	✓
ASSERTED_TIME	Semantics not specified	✓
Modified SQL syntax	✓	No
(bi) temporal PK	✓	✓
Referential integrity constraints	✓	✓

**What is Still Missing?**



## What is missing in the Standard


- ✓ Support for more than one application time dimension, which means no support for:
    - ❑ Future assertion
    - ❑ UPDATE vs. CORRECTION semantics
  - ✓ Period JOIN support
  - ✓ Period AGGREGATE support
- 



# What is missing in PostgreSQL

Of course, bitemporal!

Actually, there is no SQL-supported  
temporality



## What is Missing in Bitemporal

- ✓ Syntax extensions
- ✓ Design methodologies (also missing in the Standard)
- ✓ Explicit support for transactional time dimension, although transactional dimension is not really needed (the Standard identifies `SYSTEM_END` as the time when a record stops being `CURRENT`)

# Conclusion

# Past, Present, and Future Of Temporal Databases

- ✓ The SQL Standard provides very reasonable conservative support of Temporal tables
- ✓ PostgreSQL contains everything that is needed for efficient implementation but not an implementation
- ✓ Bitemporal is not too far from the Standard

